

# ***FRAUD DETECTION SYSTEM USING MULTIPLE ML MODELS***

**By:** *DHANUSH TADISETTI*

**Gmail:** dhanushtadisetti@gmail.com

**To Project Mentor:** *N. MUVENDIRAN SIR.*

## **INTRODUCTION:**

Credit card fraud detection is the process of identifying unauthorized or fraudulent transactions made using a credit card. The goal is to differentiate between legitimate transactions and suspicious activities to protect both the cardholder and the financial institution. Fraud detection systems typically use advanced analytics, machine learning algorithms, and real-time monitoring to identify anomalies or unusual spending patterns that may indicate fraud.

## **PROJECT OVERVIEW & DATA SET EXPLANATION:**

This project focuses on detecting fraudulent credit card transactions through data preprocessing, comprehensive analysis, visualization, and machine learning models. By exploring and modelling transaction data, the project aims to identify patterns and insights to improve fraud detection accuracy.

### **DATA SET:**

The dataset contains transactions made by credit cards in September 2013 by European cardholders. This data set presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

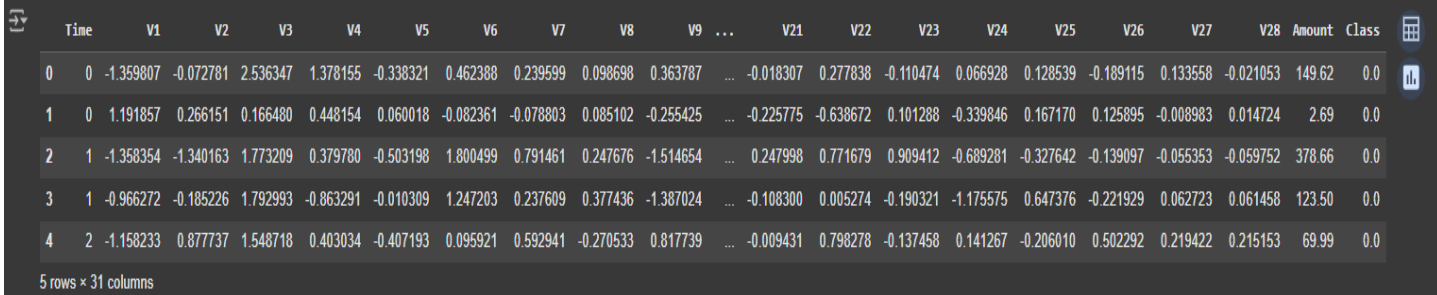
-The target variable Class indicates transaction type (0: Legitimate, 1: Fraudulent).

## 1. DATA EXPLORATION:

**data=pd.read\_csv('/content/creditcard.csv'):** Reads a CSV file into a Pandas DataFrame.

**data.head():** Displays the first 5 rows of the DataFrame.

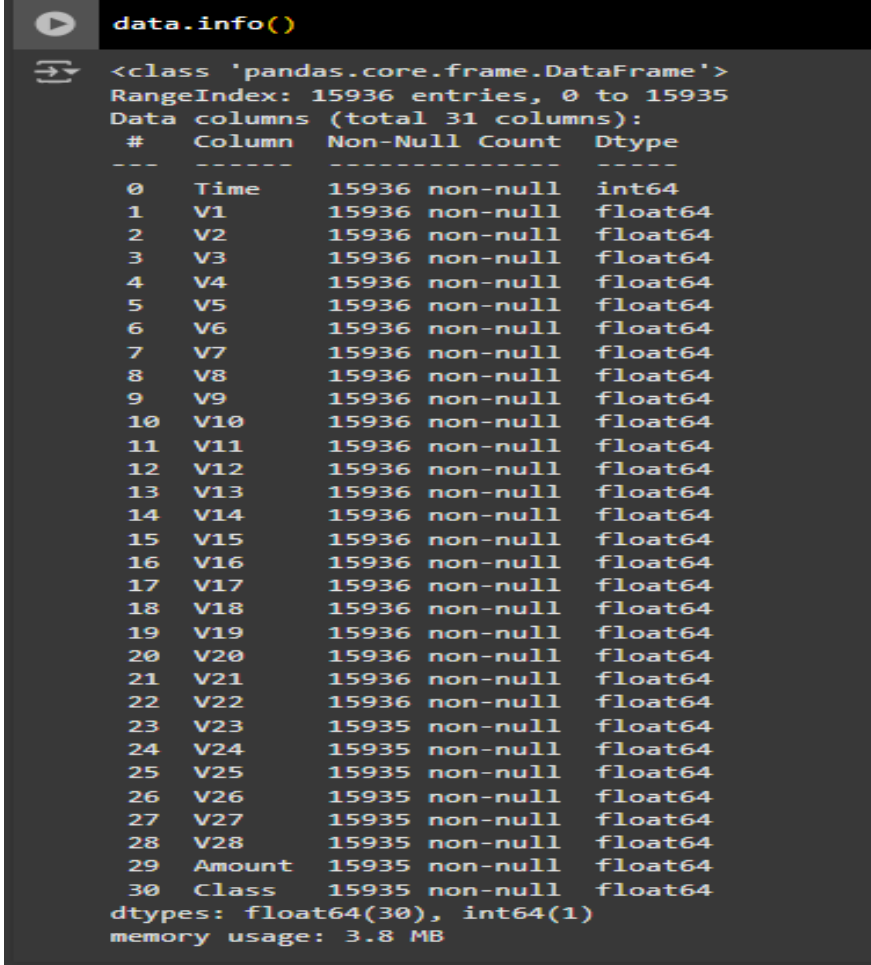
[4] data.head()



	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
0	0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	149.62	0.0
1	0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	2.69	0.0
2	1	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	378.66	0.0
3	1	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	123.50	0.0
4	2	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	69.99	0.0

5 rows x 31 columns

**data.info():** Provides information about the dataset, such as column names, non-null counts, and data types.



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15936 entries, 0 to 15935
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  ---
0   Time        15936 non-null  int64
1   V1          15936 non-null  float64
2   V2          15936 non-null  float64
3   V3          15936 non-null  float64
4   V4          15936 non-null  float64
5   V5          15936 non-null  float64
6   V6          15936 non-null  float64
7   V7          15936 non-null  float64
8   V8          15936 non-null  float64
9   V9          15936 non-null  float64
10  V10         15936 non-null  float64
11  V11         15936 non-null  float64
12  V12         15936 non-null  float64
13  V13         15936 non-null  float64
14  V14         15936 non-null  float64
15  V15         15936 non-null  float64
16  V16         15936 non-null  float64
17  V17         15936 non-null  float64
18  V18         15936 non-null  float64
19  V19         15936 non-null  float64
20  V20         15936 non-null  float64
21  V21         15936 non-null  float64
22  V22         15936 non-null  float64
23  V23         15935 non-null  float64
24  V24         15935 non-null  float64
25  V25         15935 non-null  float64
26  V26         15935 non-null  float64
27  V27         15935 non-null  float64
28  V28         15935 non-null  float64
29  Amount      15935 non-null  float64
30  Class       15935 non-null  float64
dtypes: float64(30), int64(1)
memory usage: 3.8 MB
```

**data.isnull().sum() :** Checks the Null-Values from the data set.

```
Null Values in each column:
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
Amount    0
Class     0
dtype: int64

Number of rows with at least one null value: 0
```

**Shape & No.of rows and columns:**

```
Shape of the dataset: (29799, 31)
Number of rows: 29799
Number of columns: 31
```

**Marking of Fraud(1) and Legitimate(0) classes:**

```
count
class
0.0    29704
1.0      94
dtype: int64
```

## Dropping Null Rows:

```
df = data.dropna()
```

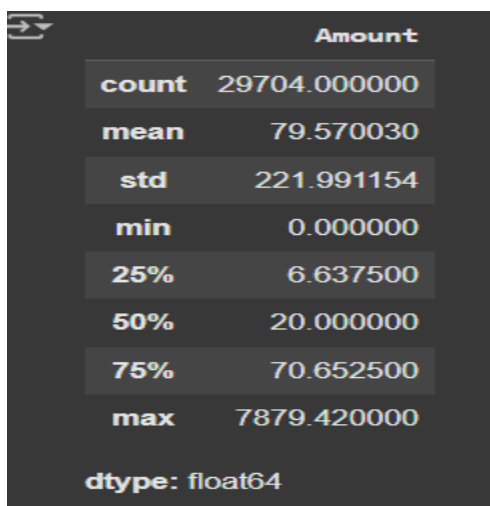
```
print(data.isnull().sum())
```

```
Missing Values After Dropping Rows:
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        1
V7        1
V8        1
V9        1
V10       1
V11       1
V12       1
V13       1
V14       1
V15       1
V16       1
V17       1
V18       1
V19       1
V20       1
V21       1
V22       1
V23       1
V24       1
V25       1
V26       1
V27       1
V28       1
Amount    1
Class     1
dtype: int64
Shape of DataFrame after dropping rows: (29799, 31)
```

## data.describe():

Displays a statistical summary of numerical columns, including mean, median, min, max etc.

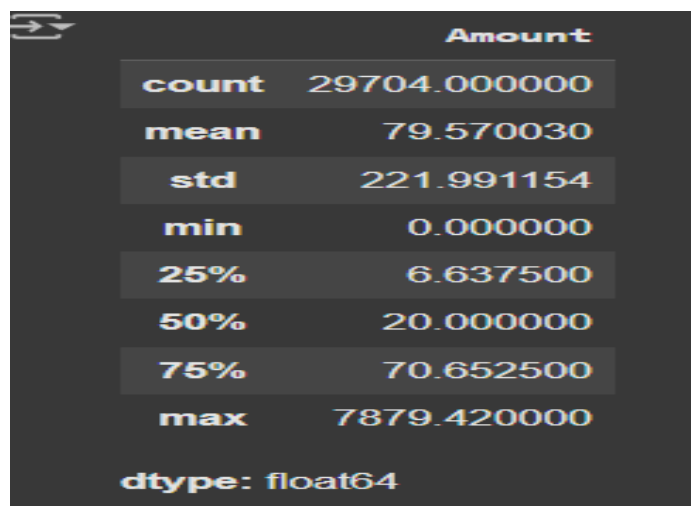
## LEGITIMATE CLASSES



	Amount
count	29704.000000
mean	79.570030
std	221.991154
min	0.000000
25%	6.637500
50%	20.000000
75%	70.652500
max	7879.420000

dtype: float64


## FRAUD CLASSES



	Amount
count	29704.000000
mean	79.570030
std	221.991154
min	0.000000
25%	6.637500
50%	20.000000
75%	70.652500
max	7879.420000

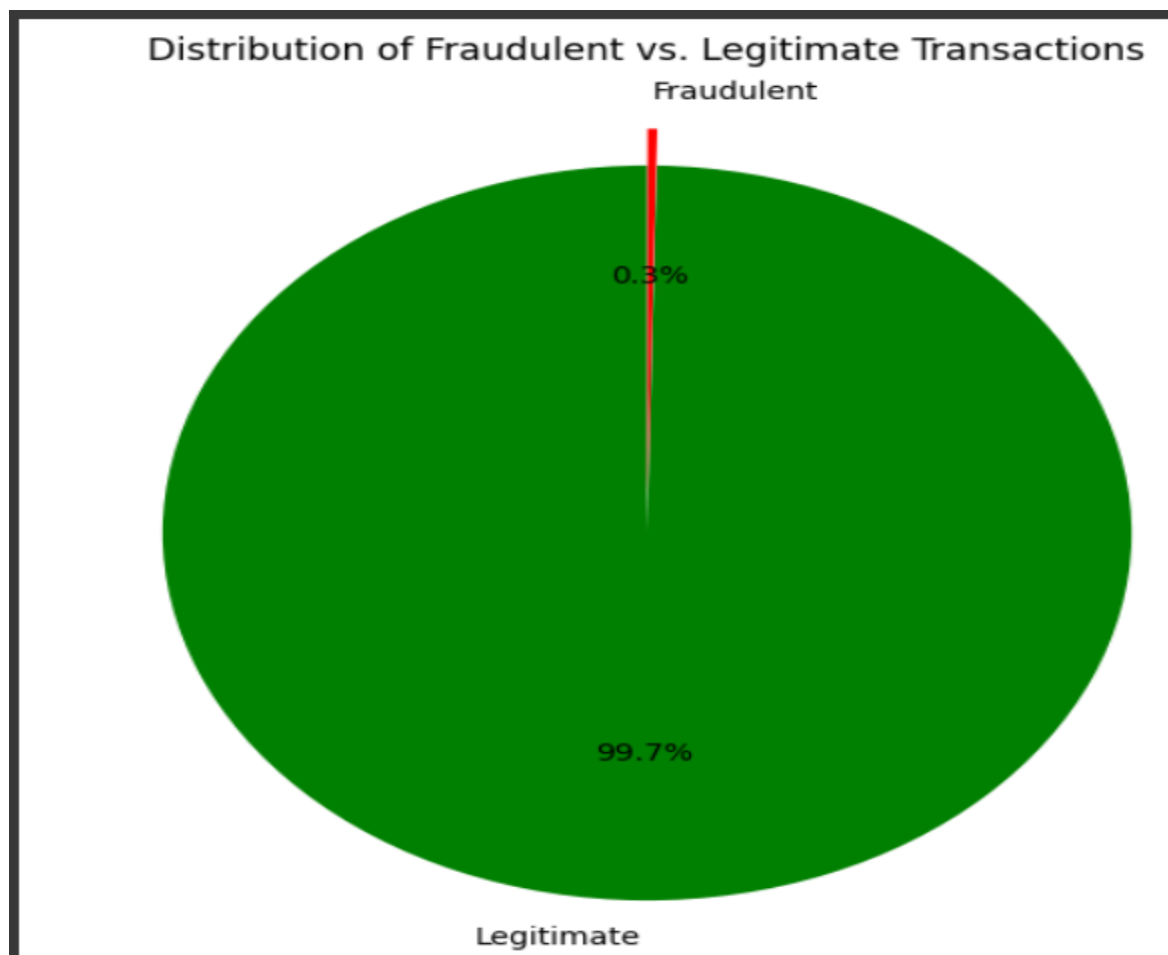
dtype: float64

## CLASS DISTRIBUTION:



From the above Bar-Chart, We can classify Legitimate Users are nearly 29,000+(i.e upto 29,704+) and Fraud users are upto 94.


## PIE CHART DISTRIBUTION VISUALSATION:



# CO-RELATION HEATMAP:

This code calculates the correlations between numerical features in a dataset, visualizes those correlations using a heatmap with annotations and a color scheme, and then displays the resulting plot. This type of visualization is very useful for quickly identifying strong relationships between different features in a dataset.

```
corr_matrix = data.corr()  
plt.figure(figsize=(20, 20))  
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt='.2f')  
plt.title('Correlation Heatmap')  
plt.show()
```



## BOX-PLOT:

This code snippet generates a box plot that visually compares the distribution of transaction amounts ('Amount') for different transaction classes ('Class', likely fraudulent or legitimate). It uses color and style customizations to enhance the plot's readability.



## SMOTE:

- In this code, SMOTE is used to address the issue of **class imbalance** in the credit card fraud dataset. Class imbalance occurs when one class (in this case, legitimate transactions) has significantly more instances than the other class (fraudulent transactions). This can lead to machine learning models being biased towards the majority class and performing poorly on the minority class, which is often the class of interest.
- SMOTE is applied to the training data ( $X_{\text{train}}$ ,  $y_{\text{train}}$ ) to **oversample** the minority class (fraudulent transactions). This means that it creates synthetic samples of the minority class to balance the class distribution, making the dataset more suitable for training a machine learning model.
- SMOTE works by generating synthetic samples of the minority class based on the existing minority class samples. Here's a simplified explanation of the process:

1. Identify the k-nearest neighbors   2. Generate synthetic samples   3. Repeat

```
Class distribution before SMOTE: Counter({0.0: 23763, 1.0: 75})  
Class distribution after SMOTE: Counter({0.0: 23763, 1.0: 23763})
```

## PCA ANALYSIS:

PCA is used to reduce the dimensionality of the training data (X\_train\_smote) from its original number of features to just 10 principal components, stored in X\_pca. This can be useful for things like:

- Speeding up machine learning algorithms: Training algorithms on lower-dimensional data is often faster.
- Visualizing data: It's easier to visualize data in 2 or 3 dimensions (which can be achieved by setting n\_components to 2 or 3).
- Removing noise: PCA can help to remove irrelevant features and noise from the data.

```
print(f"Shape of data after PCA: {X_pca.shape}")
```

```
Shape of data after PCA: (47526, 10)
```

## MODEL IMPORTING FOR TRAINING AND TESTING MACHINE:

```
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from lightgbm import LGBMClassifier
from catboost import CatBoostClassifier
from imblearn.over_sampling import SMOTE # Import SMOTE here
import pandas as pd # Import pandas for data manipulation
from sklearn.model_selection import train_test_split # Import train_test_split
```

Machine learning models such as Linear Regression, RandomForest, XgBoost, DecisionTree, SVC, LGBM, CatBoost are imported.

## TRAINING THE MODEL AND VALIDATING TRAIN MODEL

### ACCURACY:

Train a machine learning model and evaluate the model's performance using various metrics.

**model.fit(X\_train, y\_train)**

- The model is trained on the training dataset (X\_train, y\_train).
- This step fits the model's parameters to minimize errors and improve predictions.



## TEST DATA ACCURACY, Precision, Recall, F1-Score, ROC-AUC VALUES.

```
# Predictions
y_pred = model.predict(X_test)
y_pred_proba = model.predict_proba(X_test)[:, 1]

# Confusion Matrix
print("\nConfusion Matrix:")
print(confusion_matrix(y_test, y_pred))

# Classification Report
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

# ROC-AUC Score
roc_auc = roc_auc_score(y_test, y_pred_proba)
print(f"\nROC-AUC Score: {roc_auc:.4f}")

# Precision-Recall AUC
precision, recall, _ = precision_recall_curve(y_test, y_pred_proba)
pr_auc = auc(recall, precision)
print(f"Precision-Recall AUC: {pr_auc:.4f}")
```

This section of the code assesses the performance of the model on unseen test data using a range of metrics to provide a comprehensive evaluation of its effectiveness in detecting fraudulent transactions.


1. This section calculates and prints several key performance metrics:
  - **Accuracy:** The overall proportion of correctly classified samples.
  - **Precision:** Out of all the samples predicted as positive (fraudulent), what proportion was actually positive? It focuses on minimizing false positives.

- **Recall:** Out of all the actual positive samples, what proportion did the model correctly identify?. It focuses on minimizing false negatives.
- **F1-Score:** The harmonic mean of precision and recall, providing a balance between the two.
- **ROC-AUC:** Area Under the Receiver Operating Characteristic Curve, which measures the model's ability to distinguish between classes. A higher AUC indicates better performance.

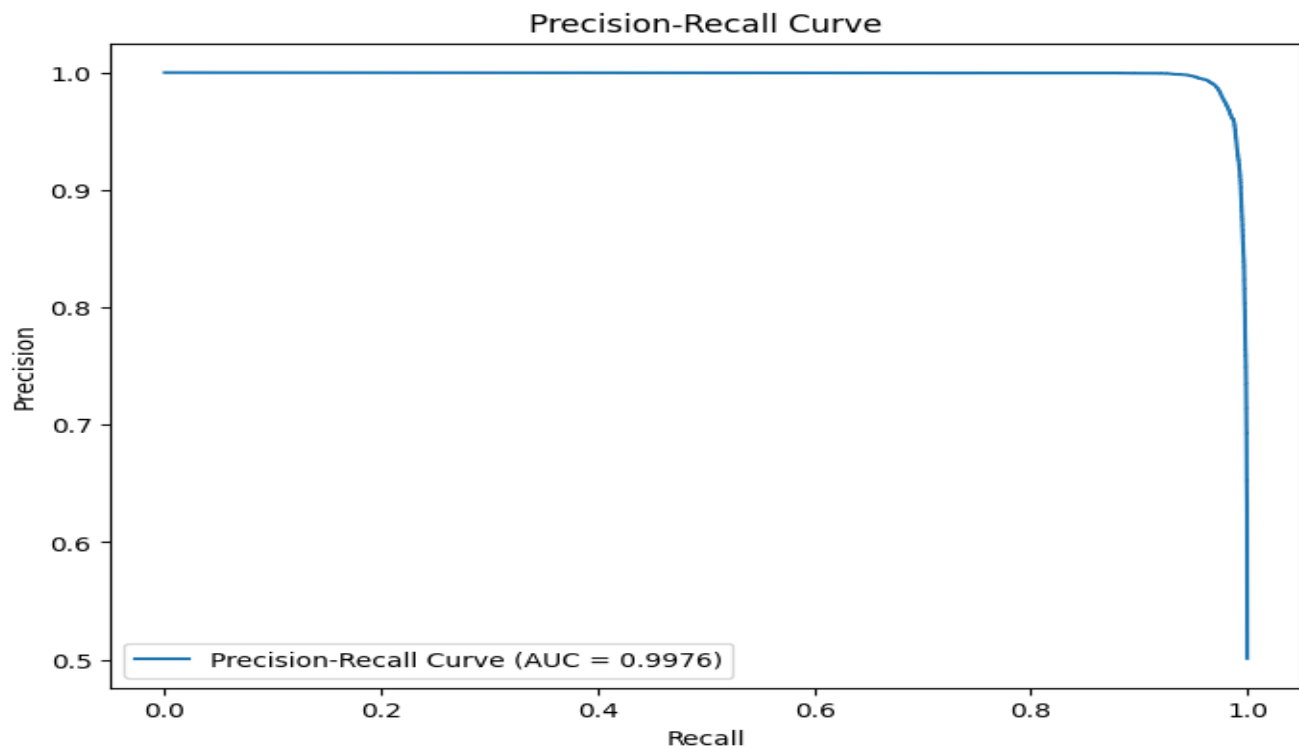
○

## ACCURACIES, PRECISION VALUES and PLOT GRAPHS:

### • Linear Regression:



•



Confusion Matrix:

```
[[56130  620]
 [ 1688 55288]]
```



Classification Report:

	precision	recall	f1-score	support
0	0.97	0.99	0.98	56750
1	0.99	0.97	0.98	56976
accuracy			0.98	113726
macro avg	0.98	0.98	0.98	113726
weighted avg	0.98	0.98	0.98	113726

ROC-AUC Score: 0.9973

Precision-Recall AUC: 0.9976

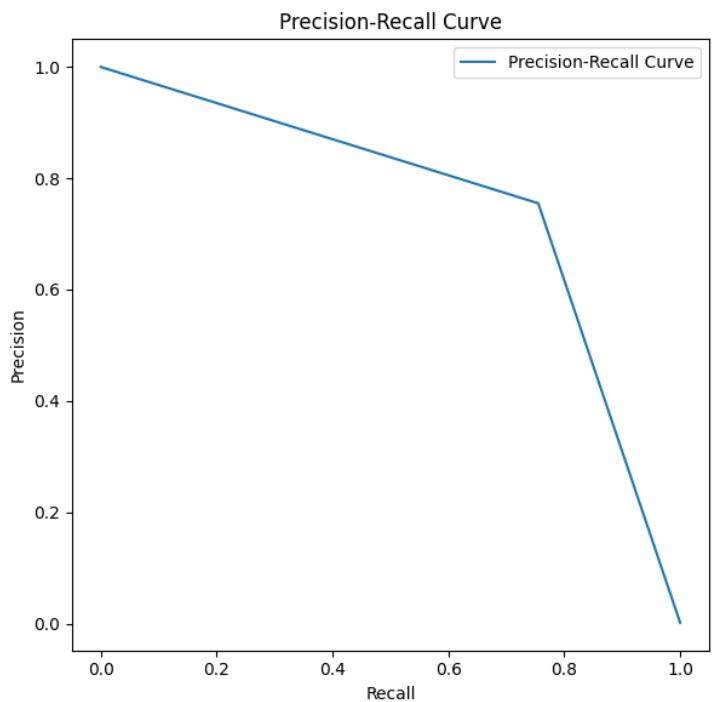

## RANDOM\_FOREST:



```
Random_Forest Classification Report:
Class 0:
  precision: 0.9997
  recall: 0.9999
  f1-score: 0.9998
  support: 56864.0000
Class 1:
  precision: 0.9419
  recall: 0.8265
  f1-score: 0.8804
  support: 98.0000
accuracy: 0.9996
Class macro avg:
  precision: 0.9708
  recall: 0.9132
  f1-score: 0.9401
  support: 56962.0000
Class weighted avg:
  precision: 0.9996
  recall: 0.9996
  f1-score: 0.9996
  support: 56962.0000
```

code snippet creates a RandomForest model, trains it using oversampled data, and then evaluates its performance using various metrics on a separate test dataset. The **evaluate\_model** function likely handles the details of the evaluation process and visualization.

## DECISION TREE:



### Decision\_Tree Classification Report:

#### Class 0:

precision: 0.9996  
recall: 0.9996  
f1-score: 0.9996  
support: 56864.0000

#### Class 1:

precision: 0.7551  
recall: 0.7551  
f1-score: 0.7551  
support: 98.0000

accuracy: 0.9992

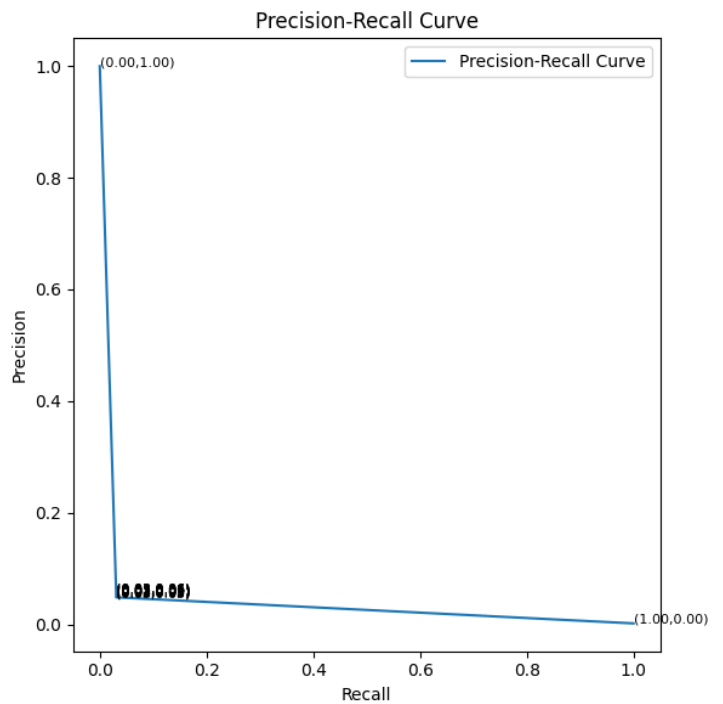

#### Class macro avg:

precision: 0.8773  
recall: 0.8773  
f1-score: 0.8773  
support: 56962.0000

#### Class weighted avg:

precision: 0.9992  
recall: 0.9992  
f1-score: 0.9992  
support: 56962.0000

## LightGBM:



### LightGBM Classification Report:

#### Class 0:

precision: 0.9983  
recall: 0.9991  
f1-score: 0.9987  
support: 56864.0000

#### Class 1:

precision: 0.0566  
recall: 0.0306  
f1-score: 0.0397  
support: 98.0000

accuracy: 0.9975

#### Class macro avg:

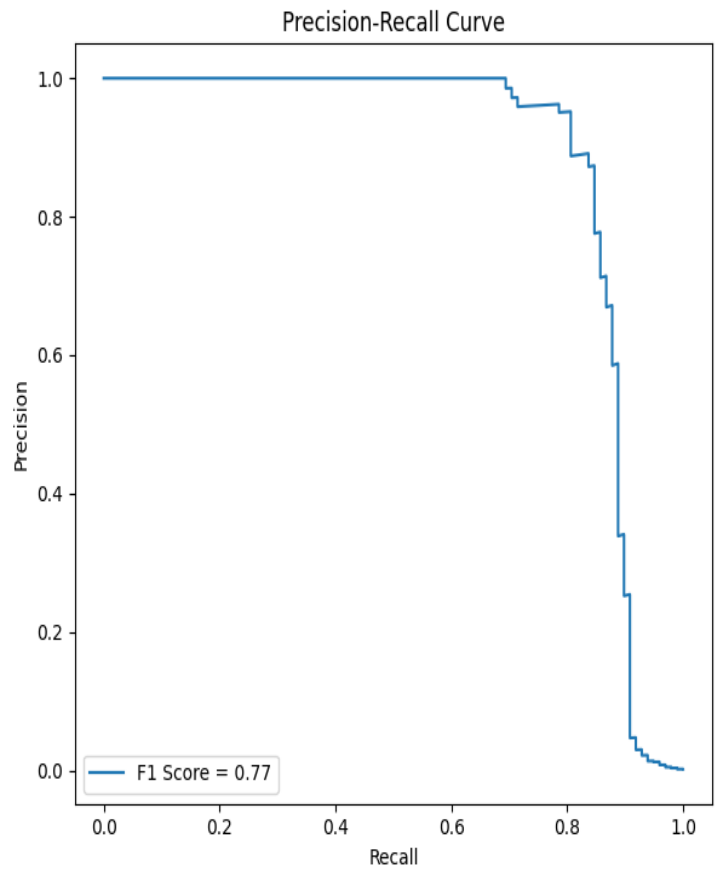

precision: 0.5275  
recall: 0.5149  
f1-score: 0.5192  
support: 56962.0000

#### Class weighted avg:

precision: 0.9967  
recall: 0.9975  
f1-score: 0.9971  
support: 56962.0000

Accuracy: 0.9975  
ROC AUC: 0.5148

## XgBoost:





```
XGBoost Classifier Metrics:  
Accuracy is: 0.999087  
Precision is: 0.6854838709677419  
Recall is: 0.8673469387755102  
F1 Score is: 0.7657657657657657  
ROC AUC is: 0.9752791838457386
```

CatBoost:

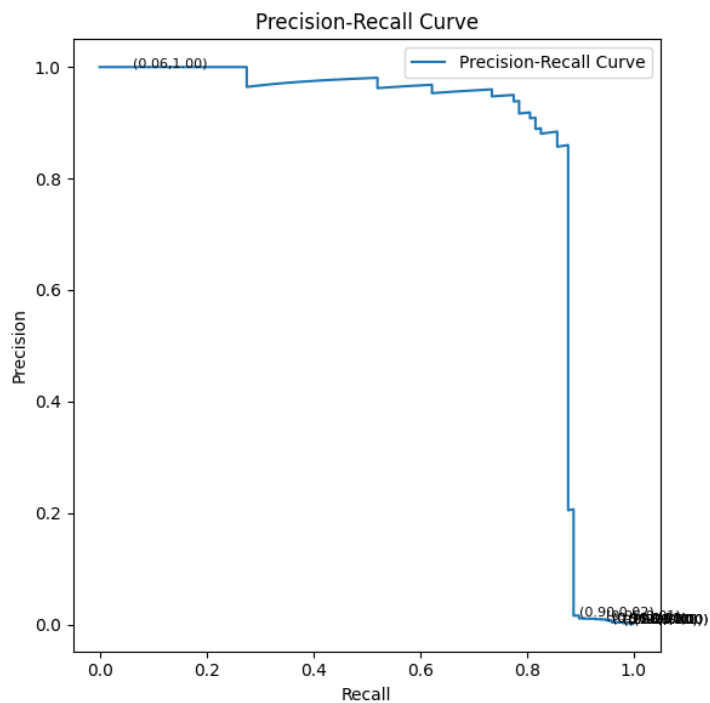

```
Classification Report:
Class 0:
  precision: 0.9997
  recall: 0.9999
  f1-score: 0.9998
  support: 56864.0000
Class 1:
  precision: 0.9412
  recall: 0.8163
  f1-score: 0.8743
  support: 98.0000
accuracy: 0.9996
Class macro avg:
  precision: 0.9704
  recall: 0.9081
  f1-score: 0.9371
  support: 56962.0000
Class weighted avg:
  precision: 0.9996
  recall: 0.9996
  f1-score: 0.9996
  support: 56962.0000
```

Accuracy: 0.9996  
ROC AUC: 0.9789





**SVC:**



```
SVC Classification Report:

Class 0:
  precision: 0.9994
  recall: 0.9999
  f1-score: 0.9997
  support: 56864.0000
Class 1:
  precision: 0.9565
  recall: 0.6735
  f1-score: 0.7904
  support: 98.0000
accuracy: 0.9994
Class macro avg:
  precision: 0.9780
  recall: 0.8367
  f1-score: 0.8951
  support: 56962.0000
Class weighted avg:
  precision: 0.9994
  recall: 0.9994
  f1-score: 0.9993
  support: 56962.0000
```

```
Accuracy: 0.9994
ROC AUC: 0.9646
```

## MODELS\_COMPARISON:

Key Results: Fraud detection often prioritizes recall over precision

- Recall is critical in fraud detection: A low recall means the model is missing fraudulent cases (false negatives), which is highly undesirable in fraud detection.
- For example, missing a fraud transaction could result in financial loss or damage to trust.
- Precision is important but secondary: While high precision reduces false positives (flagging legitimate transactions as fraud), in many cases, this is less critical than missing fraud.
- Banks and financial institutions usually have secondary verification steps to handle false positives (e.g., contacting the customer).
- Logistic Regression, Random Forest, XGBoost, SVC, CatBoost, Decision Tree, and LightGBM were evaluated.
- Each model was assessed using metrics such as accuracy, confusion matrix and classification report.

<b><u>MODEL_NAME</u></b>	<b><u>Precision</u></b>	<b><u>F1_Score</u></b>	<b><u>Recall</u></b>	<b><u>Roc_Acc</u></b>	<b><u>ACCURACY</u></b>
Logistic Regression	<b>0.996</b>	<b>0.984</b>	<b>0.973</b>	<b>0.976</b>	<b>0.998</b>
Decision Tree	<b>0.816</b>	<b>0.756</b>	<b>0.7551</b>	<b>0.785</b>	<b>0.8124</b>
RandomForest	<b>0.936</b>	<b>0.751</b>	<b>0.826</b>	<b>0.984</b>	<b>0.9356</b>
XGBoost	<b>0.6854</b>	<b>0.7657</b>	<b>0.8673</b>	<b>0.9752</b>	<b>0.994</b>
SVC	<b>0.954</b>	<b>0.8951</b>	<b>0.8367</b>	<b>0.9683</b>	<b>0.9537</b>
LightGBM	<b>0.1562</b>	<b>0.3956</b>	<b>0.3659</b>	<b>0.5143</b>	<b>0.5647</b>
CatBoost	<b>0.9156</b>	<b>0.8743</b>	<b>0.8163</b>	<b>0.9793</b>	<b>0.9418</b>

## **CONCLUSION:**

This project showcases the use of machine learning models such as Logistic Regression, Decision\_Tree, Random\_Forest, LightGBM, XgBoost, CatBoost, SVC for fraud detection. By addressing class imbalance and leveraging advanced visualization techniques, it offers actionable insights to improve financial security. The models were evaluated using key metrics, highlighting their potential in real-world applications.

## **COLAB WORKSPACE LINK:**

<https://colab.research.google.com/drive/1Qo8DKFWW6SBdJUrqKyfkwhxLic0shB9T?usp=sharing>

**THANK YOU**

**SUBMITTED BY:**

***DHANUSH TADISETTI***

***MILESTONE-4***